

Highlights

Lightweight Chinese News Summarization via Knowledge Distillation and Curriculum Learning

Yu-Hao Wang, Shang-Liang Chen

- Introduces a **five-stage curriculum learning framework** tailored for compact LLMs.
- Combines **knowledge distillation** with progressive training to improve fluency.
- Proposes a **summary-first data generation** strategy that yields coherent and focused training samples.
- Achieves **3B-tier summarization quality** with a **0.5B-parameter model**.
- Corrects **teacher-induced structural and formatting errors**, improving Traditional Chinese summary quality.

Lightweight Chinese News Summarization via Knowledge Distillation and Curriculum Learning

Yu-Hao Wang^a, Shang-Liang Chen^{a,*}

^aDepartment of Computer Science and Information Engineering, National Cheng Kung University, No. 1, University Road, Tainan, 701, Taiwan

Abstract

In today's era of information overload, readers often lack time to engage with full news articles and may be misled by sensational headlines. This work presents a lightweight Chinese news summarization model designed for resource-limited devices while maintaining strong comprehension and fluency. By combining knowledge distillation with a five-stage curriculum, covering traditional Chinese conversion, aspect extraction, reasoning construction, and summary generation, the model achieves efficient parameter compression without sacrificing quality. We further evaluate data generation methods and training strategies, including staged learning-rate adjustment, selective parameter freezing, and LoRA adaptation. Experiments show that resetting the learning rate in the final stage enhances optimization stability, while full-parameter fine-tuning remains the most effective strategy. Contrary to expectations, LoRA and partial freezing offer minimal efficiency gains. The trained 0.5B-parameter model matches or surpasses larger counterparts, generalizes well to unseen datasets, and effectively eliminates structural and formatting errors inherited from the teacher. These results demonstrate the viability of compact yet high-quality summarization models suitable for practical mobile deployment.

Keywords: chinese news summarization, lightweight models, knowledge distillation, curriculum learning

*Corresponding author.

Email addresses: f74116152@gs.ncku.edu.tw (Yu-Hao Wang),
slchen@mail.ncku.edu.tw (Shang-Liang Chen)

1. Introduction

In today’s fast-paced society, readers often lack the time to thoroughly engage with complete news articles. Sensational headlines and fragmented information can easily distort public understanding. A compact Chinese summarization model capable of running on mobile devices can address this issue by delivering concise, meaningful summaries, thereby improving both the efficiency and quality of information acquisition.

This research aims to integrate knowledge distillation and curriculum training strategies to compress model parameters while enhancing generalization and language comprehension. The ultimate objective is to develop and deploy a high-quality Chinese news summarization system on mobile devices—achieving a balance between compactness and performance.

2. Methodology

2.1. Data Collection

To build the training corpus, a custom web crawler was developed to collect news articles from *United Daily News*. Additionally, the YeungNLP/firefly-pretrain-dataset was incorporated as supplementary pretraining data. This combination provided both domain-specific news coverage and a broad linguistic foundation for model training.

2.2. Model Selection

Two models from the Qwen2.5 family were adopted [1], both based on the Transformer architecture [2]. The student model, Qwen2.5-0.5B-Instruct, is a lightweight 0.5-billion-parameter model optimized for efficient fine-tuning and deployment. The teacher model, Qwen2.5-32B-Instruct, was used for generating synthetic data, providing reasoning guidance, and offering evaluation signals during experiments.

2.3. Curriculum Training

Training was structured into five progressive stages (S1 – S5) designed to gradually increase task complexity. First, traditional Chinese text was standardized using OpenCC (S1). Next, essential aspects were extracted from each article (S2), followed by the construction of semantic triples that captured relationships among those aspects (S3), inspired by TriSum [3]. Building on this foundation, the model generated summaries conditioned

on the article content and the extracted aspects/triples (S4). Finally, the model was trained to produce summaries directly from raw news content (S5), representing the most challenging stage.

2.4. Data Generation Strategies

Several strategies for generating training data were explored. The **V1** strategy generated aspects, triples, and summaries jointly in a single step. **V2** followed a sequential pipeline—first generating aspects, then triples, and finally summaries. **V3** reversed this process by generating the summary first and then inferring supporting aspects and triples. Finally, **V4** extended V3 with manual error correction to reduce noise and improve overall quality.

2.5. Training Strategies

We explored several fine-tuning strategies to identify the most effective setup for curriculum training. The baseline employed a standard decaying learning-rate schedule, gradually lowering the rate as training progressed from Stage 1 to Stage 5. However, this monotonic decay risked premature convergence in later, more complex stages. To address this, the variant `lr_adj` retained the decay pattern in Stages 1–4 but reset the starting learning rate in Stage 5 to match the initial value of Stage 1. This adjustment provided a renewed optimization signal at the final, most demanding stage, while still allowing intra-stage decay to stabilize learning.

We also examined selective parameter freezing to analyze the contribution of specific network components: `only_attn` froze all parameters except attention layers, and `only_mlp` froze all parameters except MLP layers. Additionally, LoRA-based adaptation [4] was applied to both MLP (rank 160) and attention (rank 32) layers, introducing low-rank trainable matrices to reduce the number of fine-tuned parameters. Combinations of these strategies were further explored to identify potential synergies.

2.6. Evaluation Metrics

Model outputs were evaluated using both automatic and teacher-aligned metrics. ROUGE-1, ROUGE-2, and ROUGE-L measured lexical n-gram overlap with reference summaries, while BERTScore [5] captured semantic similarity. In addition, the teacher model (32B) acted as a judge to evaluate naturalness and information coverage, providing a complementary qualitative evaluation to automatic metrics.

Model	R-1	B-F1	Judge	R-2	R-L
4stg_v3	45.5	77.9	70.3	24.3	37.6
4stg_v1	43.8	76.8	64.0	22.1	35.5
4stg_v2	37.6	69.4	65.1	17.5	23.4

Table 1: Performance Comparison by Data Generation Strategy

3. Experiments

3.1. Comparison of Data Generation Strategies

The first experiment evaluated how different data generation strategies affected downstream summarization. Direct reasoning (V1), which required the teacher to infer aspects and triples before summarization, produced incomplete and incoherent outputs. Although some logical consistency remained, the generated summaries lacked comprehensive content coverage, reducing their usefulness for training.

The stepwise approach (V2) aimed to mitigate this issue by decomposing generation into multiple stages. However, this actually worsened coherence, producing “out-of-focus” results in which key elements deviated from the final summary. This suggests that excessive decomposition may amplify inconsistencies instead of reducing them.

In contrast, the summary-first strategy (V3) showed clear advantages. By first generating a coherent summary and then deriving supporting aspects and triples, it achieved the highest scores across all evaluation metrics (ROUGE, BERTScore, and Judge). The manually corrected V4 variant further improved data quality, establishing it as the preferred strategy.

3.2. Training Strategy Comparison

The second experiment compared the learning-rate schedules and parameter-freezing techniques described above. The `lr_adj` variant, which restored the Stage-5 starting learning rate to the initial value, achieved superior performance compared with a strictly decaying schedule. This suggests that small student models benefit from stronger optimization in the final stage, where summarization tasks are most complex.

Both `only_attn` and `only_mlp` configurations achieved competitive results, indicating that attention and MLP layers contribute comparably to summarization performance. For instance, the `only_mlp` variant maintained a high BERTScore (78.6%), demonstrating robust semantic alignment.

Model	R-1	B-F1	Judge	R-2	R-L
4stg_v3-lr_adj	48.4	79.3	72.8	25.7	40.1
4stg_v3-lr_adj-only_mlp	46.6	78.6	71.5	24.2	38.4
4stg_v3-lr_adj_lora	45.6	78.0	73.6	23.3	37.4
4stg_v3	45.5	77.9	70.3	24.3	37.6
4stg_v3-lr_adj-only_attn	45.2	77.8	69.1	23.0	37.1

Table 2: Performance Comparison by Training Strategy

Model	R-1	B-F1	Judge	R-2	R-L
Qwen2.5-0.5B_4stg_v3	45.5	77.9	70.3	24.3	37.6
Qwen2.5-0.5B_1stg_v3	45.9	77.9	68.8	23.7	37.7

Table 3: Performance of Different Stage Numbers with Decaying Learning Rate

In contrast, LoRA adaptation did not improve results in this context—performance either plateaued or slightly declined. Overall, these findings highlight that full-parameter fine-tuning with a carefully staged learning-rate reset remains the most effective strategy for compact summarization models (see Table 2).

3.3. Effect of Curriculum Stages

The third experiment examined the impact of curriculum stage count. While single-stage setups occasionally achieved slightly higher ROUGE-1 scores, multi-stage training consistently produced better linguistic quality and output stability, especially in Traditional Chinese generation.

Curriculum benefits were also found to depend on pretraining scale. When pretraining was extensive, the improvements were modest; however, with smaller pretraining datasets, curriculum learning provided substantial gains, helping compensate for weaker initial representations.

3.4. Generalization to Unseen Datasets

To evaluate generalization, inspired by the work of [11], the trained model was tested on two unseen corpora: *CLTS* [9] (long-form) and *CNewSum(v2)* [10] (headline-style). These datasets differ greatly from the training domain in style and topic diversity, providing a robust test of adaptability. Evaluation was conducted using ROUGE-1, ROUGE-2, and ROUGE-L metrics to measure cross-domain robustness.




Figure 1: Proportion of Traditional Chinese responses generated by each model. The proposed 0.5B model with 5-stage curriculum achieves the highest ratio, surpassing even 3B models.

Model	R-1	B-F1	Judge	R-2	R-L
Qwen2.5-0.5B_1stg_v4-lr_adj	47.0	78.8	75.9	24.5	38.9
Qwen2.5-0.5B_4stg_v4-lr_adj	46.8	78.6	75.2	24.2	38.6
Qwen2.5-0.5B_5stg_v4-lr_adj	46.5	78.6	75.1	23.9	38.3

Table 4: Performance of Different Stage Numbers without Decaying Learning Rate

Table 5: Performance of Custom Pretrained Models with Different Curriculum Stages

Model	R-1	B-F1	Judge	R-2	R-L
Custom-pretrained-4stg_v3-lr_adj	13.3	50.6	0.9	1.8	4.0
Custom-pretrained-1stg_v3-lr_adj	11.8	50.5	0.6	1.5	5.0

4. Results

4.1. Final Model Performance

The final Qwen2.5-0.5B model achieved performance approaching that of 3B-scale models in all automatic metrics, marking significant improvements in abstraction and accuracy. It outperformed other small-scale models such as DeepSeek [6], Gemma [7], and LLaMA [8], with ROUGE-1 gains of 7–11%. Compared with its pre-trained baseline, the post-trained model’s Judge score improved by 0.16, confirming notable gains in naturalness, informativeness, and abstractive capability.

4.2. Traditional Chinese Generation

In addition to overall performance, the model demonstrates marked improvements in linguistic handling. After post-training, the proportion of Traditional Chinese output rises significantly, even exceeding that of the teacher model. The percentage of responses completely free of Simplified Chinese characters is also higher than in other models, highlighting the effectiveness of the training data and curriculum strategy in producing culturally and linguistically appropriate text.

Beyond accuracy, the model showed substantial improvement in linguistic alignment. After post-training, the proportion of Traditional Chinese outputs rose markedly—surpassing even the teacher model. The share of responses entirely free of Simplified Chinese characters also increased, demonstrating that the curriculum design effectively enhanced cultural and linguistic fidelity.

Table 6: Performance Comparison of Final 0.5B Model with Other Models of Similar Parameter Scale

Model	R-1	B-F1	Judge	R-2	R-L
Qwen2.5-32B	55.6	82.2	81.7	34.9	48.2
Qwen2.5-3B	49.5	79.8	83.8	26.2	40.1
Qwen2.5-0.5B_5stg_v4-lr_adj	46.5	78.6	75.1	23.9	38.3
Llama-3.2-3B	43.3	76.8	74.3	20.7	34.7
Gemma-3-1B	41.8	76.0	72.6	18.3	31.2
Qwen2.5-0.5B	39.7	74.8	58.7	18.5	30.9
Gemma-2-2B	39.3	71.9	83.4	18.3	26.2
Llama-3.2-1B	38.7	74.1	63.4	16.9	29.5
DeepSeek-R1-1.5B	32.2	71.3	65.8	12.0	21.4

4.3. Generalization to Unseen Datasets

The trained 0.5B model generalized well to unseen datasets, consistently outperforming its untrained baseline and approaching 3B-tier models in all ROUGE metrics on both *CLTS* and *CNewSum(v2)* (Table 7). Although not surpassing the 32B teacher, its proximity to 3B performance highlights strong abstraction and factual alignment despite its compact size.

These results confirm that the proposed curriculum and fine-tuning strategies enable compact models to achieve an effective trade-off between efficiency and generalization, acquiring meaningful semantic representations rather than relying on memorization.

4.4. Mitigation of Teacher Model Errors

Finally, the post-trained model effectively mitigated errors inherited from the teacher. Among 30,296 samples, the teacher model produced 2,871 instances containing irregular endings or other apparent formatting errors (e.g., “\n tworzyć \n”), excluding more complex issues such as imprecise phrasing or full-/half-width inconsistencies. After training, the student model eliminated all such anomalies, confirming that staged curriculum learning and refined datasets successfully corrected structural and formatting artifacts present in the teacher outputs.

5. Conclusion

This study systematically explored strategies for developing a lightweight Chinese news summarization model, integrating data generation design, train-

Table 7: **Generalization Performance on Unseen Datasets (*CLTS* and *CNewSum(v2)*)**. Models are evaluated using ROUGE-1 (R-1), ROUGE-2 (R-2), and ROUGE-L (R-L) scores. The two highest scores in each column are shown in **bold**. Qwen2.5-32B achieves the best performance on CLTS, while Qwen2.5-3B slightly surpasses it on CNewSum(v2), demonstrating stronger domain adaptability. The proposed multi-stage trained student (Qwen2.5-0.5B_5stg_v4-lr_adj) maintains competitive generalization despite its smaller scale.

Model	CLTS			CNewSum(v2)		
	R-1	R-2	R-L	R-1	R-2	R-L
Qwen2.5-32B	25.31	14.62	24.92	19.64	4.94	19.13
Qwen2.5-3B	22.36	10.96	21.63	20.09	5.15	19.61
Qwen2.5-0.5B_5stg_v4-lr_adj	20.50	9.63	19.89	17.64	4.21	17.19
Qwen2.5-0.5B	18.61	9.03	18.25	15.63	3.73	15.28
Gemma-3-1B	17.79	7.94	17.26	16.44	3.87	16.02
Llama-3.2-3B	17.10	7.62	16.69	16.36	3.81	16.02
DeepSeek-R1-1.5B	16.45	8.63	15.96	11.10	2.51	10.72
Llama-3.2-1B	16.19	7.33	15.83	13.33	3.00	12.94
Gemma-2-2B	12.83	5.79	12.44	9.97	2.35	9.80

ing configurations, and curriculum learning. The findings offer both methodological insight and practical guidance for achieving high-quality summarization under constrained computational resources.

In terms of data generation, the experiments revealed that direct reasoning and stage-by-stage decomposition tend to produce incomplete or unfocused outputs, limiting their utility as training data. By contrast, the summary-first approach consistently yielded coherent and comprehensive summaries, providing a stable foundation for downstream reasoning and facilitating more effective student-teacher distillation.

Regarding training strategies, applying a decaying learning rate across the early stages and restoring a stronger starting rate in the final stage (`lr_adj`) yielded the best results. Combined with full-parameter fine-tuning, this approach outperformed LoRA-based and partial-freezing methods, showing that compact models benefit most from steady, end-stage optimization across all parameters. Moreover, contrary to expectations, LoRA and partial freezing provided little benefit in terms of training time or VRAM savings, reaffirming the practicality of full fine-tuning for small-scale summarization models.

The investigation of curriculum learning further emphasized its value, particularly under limited pretraining conditions. While multi-stage curric-

ula produced modest improvements when large-scale pretraining was available, they substantially enhanced summarization accuracy, stylistic fidelity, and robustness for smaller, custom-pretrained models. This progressive task structure also enabled the student model to correct structural and linguistic imperfections inherited from the teacher model, such as irregular endings and formatting inconsistencies, demonstrating the corrective power of staged learning.

Evaluation on unseen datasets (*CLTS* and *CNewSum(v2)*) confirmed the generalization capability of the proposed approach. The trained 0.5B model consistently outperformed its untrained baseline and approached the performance of 3B-tier models across ROUGE and BERTScore metrics, showing strong abstraction and robustness despite its compact scale.

Collectively, these strategies produced a 0.5B-parameter model that achieves performance levels comparable to much larger models while maintaining superior efficiency and linguistic precision. Notably, the model surpasses even its teacher in Traditional Chinese generation quality, reflecting the success of the designed data and training pipeline.

In conclusion, this work demonstrates that through summary-first data generation, steady full-parameter fine-tuning, and staged curriculum progression, a compact model can achieve a balance between efficiency, robustness, and high-quality abstractive summarization. The proposed framework not only advances the practical deployment of summarization systems on mobile devices but also provides a scalable and reproducible path toward resource-efficient natural language processing.

Appendix A. Prompt Templates and Examples

This appendix provides the trained student model, the prompt templates and examples used in this study, including the complete prompt designs for final versions of essential aspect extraction, triple generation, summary generation, and model evaluation.

Published model:

https://huggingface.co/BennyWang/Qwen2.5-0.5B-Instruct-Curriculum-5stage-v4-lr_adj

Example news, essential aspects, triples, and summary (V3):

News:

母愛不分物種。動保組織接獲民眾通報發現草叢有一窩胖嘟嘟奶汪，沒想到牠們

的母親為了照顧這些孩子把自己餓成皮包骨，還有嚴重營養不良跟脫水狀況，對比起小狗們都肥碩健康，狗媽媽更是令人心疼。

根據 The Dodo 報導，美國密蘇里州 (Missouri) 聖路易斯流浪動物救援組織 (Stray Rescue of St. Louis,SRSL) 日前接獲民眾通報，說草叢裡面發現一窩肥胖的奶汪，但卻找不到狗媽媽，希望他們可以派人來協助一下。動物救援組織工作人員湯姆森 (Natalie Thomson) 表示，當他們趕往民眾通報的現場，的確真的看到一窩被照顧得好好的奶汪，很像是被人飼養後遺棄在附近。

但令動保人員意外的一幕出現了，他們過沒多久在附近的草叢找到了狗媽媽，可是這隻渾身骨瘦如柴、幾乎可以用皮包骨形容的黃狗顯然非常營養不良，跟牠一窩肥壯幼崽形成強烈對比，這隻狗媽媽顯得有些害怕人類，但牠並沒有逃跑或是圖保護孩子，反而是將目光投射在小狗身上，希望眼前的人類不要傷害牠的孩子。每一隻狗寶寶都相當健康可愛甚至還有點肥。(圖/取自 Stray Rescue of St. Louis 官網)

後來這一窩共 10 隻奶汪跟牠們孱弱的母親都被聖路易斯流浪動物救援組織帶回收容所，經過健康檢查後反而讓獸醫跟動保人員更難過了，因為這 10 隻小奶狗除了有點寄生蟲問題之外，沒有任何營養不良的狀況，甚至還有些過重，但狗媽媽卻嚴重營養不良還脫水，可以說為了照顧孩子鞠躬盡瘁。

目前這一窩小奶汪都受到良好的照顧，不日將可開放認養，而狗媽媽則因為身體虛弱還需要靜養一段時間才可以考慮出養。湯姆森說，這隻狗媽媽其實是很溫柔的，只是個性比較慢熟，所以需要有愛心跟耐心的飼主陪伴，牠就會慢慢敞開心房願意相信人類。

Summary:

美國密蘇里州一窩小狗被發現肥胖健康，但牠們的母親卻因照顧孩子而嚴重營養不良、脫水。這 11 隻狗已被救出，小狗們即將開放認養，母親則需繼續靜養。救援組織尋找有耐心的飼主，以陪伴這位溫柔但慢熟的母親。

Essential aspects:

[美國密蘇里州]、[一窩小狗]、[肥胖健康]、[母親嚴重營養不良]、[脫水]、[已被救出]、[小狗們即將開放認養]、[母親需繼續靜養]、[救援組織尋找有耐心的飼主]、[溫柔但慢熟的母親]

Triples :

[美國密蘇里州 | 發現 | 一窩小狗], [小狗 | 肥胖健康 | 狗媽媽], [狗媽媽 | 照顧孩子 | 嚴重營養不良], [狗媽媽 | 照顧孩子 | 脫水], [一窩小狗 | 被 | 救出], [小狗們 | 即將 | 開放認養], [狗媽媽 | 需 | 繼續靜養], [救援組織 | 尋找 | 有耐心的飼主], [狗媽媽 | 是 | 溫柔但慢熟的]

Prompt for generating V3 summary:

System prompt:

請根據以下新聞內容，為新聞生成一份 100 字內精簡的摘要，請用繁體中文回答。

例如：

生成摘要：

=====

User prompt:

新聞：

{news}

Prompt for generating V3 essential aspects and triples:

System prompt:

請根據以下新聞內容以及摘要，提取新聞的關鍵要素與三元組，關鍵要素應為關鍵短句、名詞或事實，三元組應為 [實體 1 | 關係 | 實體 2] 的格式，這些三元組用於構成摘要，請用繁體中文回答。請將每個關鍵要素與三元組用 [] 與、分隔。例如：

關鍵要素：

[關鍵要素 1]、[關鍵要素 2]、[關鍵要素 3]、...

三元組：

[實體 1_1 | 關係 _1 | 實體 1_2]、[實體 2_1 | 關係 _2 | 實體 2_2]、...

=====

User prompt:

新聞：

{news}

摘要：

{summary}

Prompt for model scoring:

System prompt:

你是一位語言評估專家。你的任務是根據文章與標準摘要，評估模型生成的摘要品質。

請根據以下評分標準，從 0 到 20 為其打分：

- 0：格式不正確或無意義的文字。
- 1：完全無關，與文章毫不相干。
- 2：虛構內容，語意不明。
- 3：嚴重誤解，包含重大錯誤。
- 4：幾乎無法反映原文，非常不完整。
- 5：文法錯誤，缺乏連貫性與相關性。
- 6：內容不完整且部分離題。
- 7：遺漏關鍵要點，有輕微虛構。
- 8：摘要過於模糊，缺乏具體性。
- 9：簡潔，涵蓋大部分重點。
- 10：可理解但可能遺漏細節。
- 11：忠實但略有遺漏。
- 12：大致正確但稍顯冗餘。
- 13：準確、結構良好，但有輕微風格問題。
- 14：涵蓋完整、清晰，語氣尚可改進。
- 15：清楚、忠實且具風格。
- 16：簡潔優雅，涵蓋所有重點。
- 17：非常接近理想摘要，僅有些微瑕疵。
- 18：優秀的摘要，易讀且內容完整。
- 19：幾近完美，僅可做細微風格潤飾。
- 20：完美——清楚、忠實、完整且優雅。

請回傳”分數：”加一個整數分數（0-20），接著是一句簡短的理由（例如：「分數：17 —— 非常接近理想摘要，僅有些微瑕疵」）。

=====

User prompt:

文章：

{article}

標準摘要：

{ground_truth}

模型生成摘要：

{response}

References

- [1] Yang, A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Li, C., Liu, D., Huang, F., Wei, H., Lin, H., Yang, J., Tu, J., Zhang, J., Yang, J.,

- Yang, J., Zhou, J., Lin, J., Dang, K., Lu, K., Bao, K., Yang, K., Yu, L., Li, M., Xue, M., Zhang, P., Zhu, Q., Men, R., Lin, R., Li, T., Tang, T., Xia, T., Ren, X., Ren, X., Fan, Y., Su, Y., Zhang, Y., Wan, Y., Liu, Y., Cui, Z., Zhang, Z., Qiu, Z.: Qwen2.5 Technical Report. arXiv preprint arXiv:2412.15115 (2024). <https://doi.org/10.48550/arXiv.2412.15115>
- [2] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, Ł., Polosukhin, I.: Attention Is All You Need. In: Advances in Neural Information Processing Systems (NeurIPS), Vol. 30. Curran Associates, Inc. (2017). <https://doi.org/10.48550/arXiv.1706.03762>
- [3] Jiang, P., Xiao, C., Wang, Z., Bhatia, P., Sun, J., Han, J.: TriSum: Learning Summarization Ability from Large Language Models with Structured Rationale. In: Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL), pp. 2805–2819. (2024). <https://doi.org/10.48550/arXiv.2403.10351>
- [4] Hu, E.J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W.: LoRA: Low-Rank Adaptation of Large Language Models. arXiv preprint arXiv:2106.09685 (2021). <https://doi.org/10.48550/arXiv.2106.09685>
- [5] Zhang, T., Kishore, V., Wu, F., Weinberger, K.Q., Artzi, Y.: BERTScore: Evaluating Text Generation with BERT. In: International Conference on Learning Representations (ICLR). (2020). <https://doi.org/10.48550/arXiv.1904.09675>
- [6] DeepSeek-AI, Guo, D., Yang, D., Zhang, H., Song, J., Zhang, R., Xu, R., Zhu, Q., Ma, S., Wang, P., Bi, X., Zhang, X., Yu, X., Wu, Y., Wu, Z.F., Gou, Z., Shao, Z., Li, Z., Gao, Z., Liu, A., Xue, B., Wang, B., Wu, B., Feng, B., Lu, C., Zhao, C., Deng, C., Zhang, C., Ruan, C., Dai, D., Chen, D., Ji, D., Li, E., Lin, F., Dai, F., Luo, F., Hao, G., Chen, G., Li, G., Zhang, H., Bao, H., Xu, H., Wang, H., Ding, H., Xin, H., Gao, H., Qu, H., Li, H., Guo, J., Li, J., Wang, J., Chen, J., Yuan, J., Qiu, J., Li, J., Cai, J.L., Ni, J., Liang, J., Chen, J., Dong, K., Hu, K., Gao, K., Guan, K., Huang, K., Yu, K., Wang, L., Zhang, L., Zhao, L., Wang, L., Zhang, L., Xu, L., Xia, L., Zhang, M., Zhang, M., Tang,

M., Li, M., Wang, M., Li, M., Tian, N., Huang, P., Zhang, P., Wang, Q., Chen, Q., Du, Q., Ge, R., Zhang, R., Pan, R., Wang, R., Chen, R.J., Jin, R.L., Chen, R., Lu, S., Zhou, S., Chen, S., Ye, S., Wang, S., Yu, S., Zhou, S., Pan, S., Li, S.S. et al.: DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. (2025). <https://doi.org/10.48550/arXiv.2501.12948>

- [7] Riviere, M., Pathak, S., Sessa, P.G., Hardin, C., Bhupatiraju, S., Hussenot, L., Mesnard, T., Shahriari, B., Ramé, A., Ferret, J., Liu, P., Tafti, P., Friesen, A., Casbon, M., Ramos, S., Kumar, R., Le Lan, C., Jerome, S., Tsitsulin, A., Vieillard, N., Stanczyk, P., Girgin, S., Momchev, N., Hoffman, M., Thakoor, S., Grill, J.-B., Neyshabur, B., Bachem, O., Walton, A., Severyn, A., Parrish, A., Ahmad, A., Hutchison, A., Abdagic, A., Carl, A., Shen, A., Brock, A., Coenen, A., Laforge, A., Paterson, A., Bastian, B., Piot, B., Wu, B., Royal, B., Chen, C., Kumar, C., Perry, C., Welty, C., Choquette-Choo, C.A., Sinopalnikov, D., Weinberger, D., Vijaykumar, D., Rogozińska, D., Herbison, D., Bandy, E., Wang, E., Noland, E., Moreira, E., Senter, E., Eltyshev, E., Visin, F., Rasskin, G., Wei, G., Cameron, G., Martins, G., Hashemi, H., Klimczak-Plucińska, H., Batra, H., Dhand, H., Nardini, I., Mein, J., Zhou, J., Svensson, J., Stanway, J., Chan, J., Zhou, J.P., Carrasqueira, J., Iljazi, J., Becker, J., Fernandez, J., van Amersfoort, J., Gordon, J., Lipschultz, J., Newlan, J., Ji, J., Mohamed, K., Badola, K., Black, K., Millican, K., McDonell, K., Nguyen, K., Sodhia, K., Greene, K., Sjoesund, L.L., Usui, L., Sifre, L., Heuermann, L., Lago, L., McNealus, L. et al.: Gemma 2: Improving Open Language Models at a Practical Size. (2024). <https://doi.org/10.48550/arXiv.2408.00118>
- [8] Grattafiori, A., Dubey, A., Jauhri, A., Pandey, A., Kadian, A., Al-Dahle, A., Letman, A., Mathur, A., Schelten, A., Vaughan, A., Yang, A., Fan, A., Goyal, A., Hartshorn, A., Yang, A., Mitra, A., Sravankumar, A., Korenev, A., Hinsvark, A., Rao, A., Zhang, A., Rodriguez, A., Gregerson, A., Spataru, A., Roziere, B., Biron, B., Tang, B., Chern, B., Caucheteux, C., Nayak, C., Bi, C., Marra, C., McConnell, C., Keller, C., Touret, C., Wu, C., Wong, C., Canton Ferrer, C., Nikolaidis, C., Allonsius, D., Song, D., Pintz, D., Livshits, D., Wyatt, D., Esiobu, D., Choudhary, D., Mahajan, D., Garcia-Olano, D., Perino, D., Hupkes, D., Lakomkin, E., AlBadawy, E., Lobanova, E., Dinan, E., Smith,

- E.M., Radenovic, F., Guzmán, F., Zhang, F., Synnaeve, G., Lee, G., Anderson, G.L., Thattai, G., Nail, G., Mialon, G., Pang, G., Cucurell, G., Nguyen, H., Korevaar, H., Xu, H., Touvron, H., Zarov, I., Ibarra, I.A., Kloumann, I., Misra, I., Evtimov, I., Zhang, J., Copet, J., Lee, J., Geffert, J., Vranes, J., Park, J., Mahadeokar, J., Shah, J., van der Linde, J., Billock, J., Hong, J., Lee, J., Fu, J., Chi, J., Huang, J., Liu, J., Yu, J., Bitton, J., Spisak, J., Park, J., Rocca, J., Johnston, J., Saxe, J., Jia, J. et al.: The Llama 3 Herd of Models. (2024). <https://doi.org/10.48550/arXiv.2407.21783>
- [9] Liu, X., Zhang, C., Chen, X., Cao, Y., Li, J.: CLTS: A New Chinese Long Text Summarization Dataset. In: Proceedings of the 28th International Conference on Computational Linguistics, pp. 531–542. International Committee on Computational Linguistics (2020). https://doi.org/10.1007/978-3-030-60450-9_42
- [10] Wang, D., Chen, J., Wu, X., Zhou, H., Li, L.: CNewSum: A Large-scale Chinese News Summarization Dataset with Human-annotated Adequacy and Deducibility Level. (2021). <https://doi.org/10.48550/arXiv.2110.10874>
- [11] Basyal, L., Sanghvi, M.: Text Summarization Using Large Language Models: A Comparative Study of MPT-7b-instruct, Falcon-7b-instruct, and OpenAI Chat-GPT Models. (2023). <https://doi.org/10.48550/arXiv.2310.10449>